

Loan Approval Prediction - Executive Summary

Dataset contains 614 records. Objective is to predict loan approval using supervised machine learning. Data preprocessing includes missing value handling, categorical encoding and scaling. Logistic Regression and Tree-based models are compared using Precision, Recall, F1-score and ROC-AUC. Class imbalance is addressed using class weights or SMOTE. Recommended deployment threshold should prioritize recall while maintaining acceptable precision.